

Name: Jude Onyia
 Student ID: V00947095
 Email: judeonyia10@gmail.com
 Course: ECE596C
 Section: T01

Assignment ID: cpp_concurrency
 Assignment Title: Concurrency

Submission Source: https://github.com/uvic-seng475-2020-05/cpp_concurrency-JudeOnyia.git

Commit ID: 223db394dec9c37c53f3285049e145999463bd1a

Submitted Files

=====

```
drwxrwxr-x      98 2020-07-29 00:12 ./app
-rw-rw-r--    1116 2020-07-29 00:12 ./app/test_julia_set.cpp
-rw-rw-r--    1993 2020-07-29 00:12 ./app/test_queue.cpp
-rw-rw-r--     498 2020-07-29 00:12 ./app/test_thread_pool.cpp
-rw-rw-r--     875 2020-07-29 00:12 ./CMakeLists.txt
-rw-rw-r--     145 2020-07-29 00:12 ./IDENTIFICATION.txt
drwxrwxr-x      24 2020-07-29 00:12 ./include
drwxrwxr-x      83 2020-07-29 00:12 ./include/ra
-rw-rw-r--    2017 2020-07-29 00:12 ./include/ra/julia_set.hpp
-rw-rw-r--    6000 2020-07-29 00:12 ./include/ra/queue.hpp
-rw-rw-r--    5012 2020-07-29 00:12 ./include/ra/thread_pool.hpp
drwxrwxr-x      37 2020-07-29 00:12 ./lib
-rw-rw-r--     561 2020-07-29 00:12 ./lib/thread_pool.cpp
-rw-rw-r--   887689 2020-07-29 00:12 ./README.pdf
```

Results

=====

Package	Operation	Target	Status
nonprog	generate	---	OK (0.1s)
tpool_orig	generate	---	OK (0.5s)
tpool_orig	configure	---	OK (4.0s)
tpool_orig	build	test_queue	OK (1.3s)
tpool_orig	build	test_thread_pool	OK (1.7s)
tpool_sane	generate	---	OK (0.4s)
tpool_sane	configure	---	OK (3.8s)
tpool_sane	build	test_queue	OK (1.2s)
tpool_sane	build	test_thread_pool	OK (1.9s)
fractal_orig	generate	---	OK (0.4s)
fractal_orig	configure	---	OK (3.7s)
fractal_orig	build	test_julia_set	OK (3.8s)
fractal_sane	generate	---	OK (0.7s)
fractal_sane	configure	---	OK (3.8s)
fractal_sane	build	test_julia_set	OK (3.3s)

Normally, an operation is indicated as having a status of either "OK" or "FAIL". A status of "?" indicates that the operation could not be performed for some reason (e.g., due to an earlier error or being a manual step). The time (in seconds) required for an operation is denoted by an expression consisting of a number followed by the letter "s" (e.g., "5.0s"). In the case of a test that consists of multiple test cases, the number of failed test cases and total number of test cases is expressed as a fraction (e.g., "10/50" means 10 test cases failed out of 50 test cases in total).

The length (in lines) of the log file generated by an operation is denoted by an expression consisting of a number followed by the letter "L" (e.g., "10L"). To ascertain the reason for the failure of an operation, check the contents of the log file provided.

Legend

=====

Package: nonprog
Nonprogramming exercises

Package: tpool_orig
The code as originally submitted by the student.
Build target: test_queue
Build the test_queue program.
Build target: test_thread_pool
Build the test_thread_pool program.

Package: tpool_sane
Code with modifications to perform API sanity checking.
Build target: test_queue
Build the (dummy) test_queue program.
Build target: test_thread_pool
Build the (dummy) test_thread_pool program.

Package: fractal_orig
The code as originally submitted by the student.
Build target: test_julia_set
Build the test_julia_set program.

Package: fractal_sane
Code with modifications to perform API sanity checking.
Build target: test_julia_set
Build the (dummy) test_julia_set program.

```
1  commit 19d2b0001c79106add9efa3ecb5097509fe8653c
2  Author: JudeOnyia <60678029+JudeOnyia@users.noreply.github.com>
3  Date:   Mon Jul 27 00:54:28 2020 -0700
4
5      First commit
6
7  commit f6e196dc3a5c67e55ddca3110ef2562ce1b8ca9d
8  Author: JudeOnyia <60678029+JudeOnyia@users.noreply.github.com>
9  Date:   Mon Jul 27 12:19:51 2020 -0700
10
11     Added Sanitizers
12
13  commit 3f7ae0e44d6a141f9ef7fda32677fbbd4a243a85
14  Author: root <judeonyia10@gmail.com>
15  Date:   Mon Jul 27 19:42:58 2020 -0700
16
17     1) Wrote function to push and pop elements in and from the queue
18     2) Wrote function to clear and close the queue
19     3) Wrote is_full, is_empty, is_closed and max_size functions
20
21  commit 023de44ae5f2f842985e89a2a34cf51390db511c
22  Author: JudeOnyia <60678029+JudeOnyia@users.noreply.github.com>
23  Date:   Mon Jul 27 20:24:48 2020 -0700
24
25     1) Set up made for thread pool class
26
27  commit dcaa5f5de461c99086ecb690acfc88cccc1cc6186
28  Author: root <judeonyia10@gmail.com>
29  Date:   Tue Jul 28 02:28:57 2020 -0700
30
31     1) Prevented some potential infinite waits in queue class
32     2) Wrote the constructors and destructors of thread_pool class
33     3) Wrote the friend function Worker to be used by threads
34     4) Wrote the schedule function to put tasks on the queue
35     5) Wrote the shutdown function
36     6) Wrote the size and is_shutdown functions
37
38  commit 83961c676fb7b28c3eae3f0880e4a171919a9ce1
39  Author: root <judeonyia10@gmail.com>
40  Date:   Tue Jul 28 13:51:54 2020 -0700
41
42     1) Checks if a thread is joinable before joining in threadpool
43     destructor
44     2) Made shutdown function blocking instead of schedule function
45
46  commit 61c2b12e95fe6fb755cf31ff04e1a802e9945d5d
47  Author: JudeOnyia <60678029+JudeOnyia@users.noreply.github.com>
48  Date:   Tue Jul 28 14:06:38 2020 -0700
49
50     1) Set up for julia set
51
52  commit c1542b5b17394d01666d5a5f407176a68cfda637
53  Author: root <judeonyia10@gmail.com>
54  Date:   Tue Jul 28 20:59:57 2020 -0700
55
56     1) Wrote helper function getZ to calculate z from the info given
57     2) Wrote helper function julia_ym to calculate the ym term given z
58     3) Wrote function to compute julia set
59     4) Wrote test for function and visualized result
60
61  commit baa6932f2bbb41311db54bfff27f6b5e5488d82f9
62  Author: root <judeonyia10@gmail.com>
```

```
63 Date:    Tue Jul 28 23:25:48 2020 -0700
64
65     1) Used the chrono library to record execution time of julia set program
66
67 commit 223db394dec9c37c53f3285049e145999463bd1a
68 Author: JudeOnyia <60678029+JudeOnyia@users.noreply.github.com>
69 Date:    Wed Jul 29 00:01:50 2020 -0700
70
71     1) Recorded the execution time of julia set in README.pdf
72     2) Removed Sanitizer.cmake
```

Name: Jude Onyia
Student ID: V00947095
Course: ECE 596C
Due Date: July 19, 2020

Assignment 6: Non – Programming Exercise

7.1 a)

This program creates two threads with shared variables x and y initially set to zero. Thread $t1$ assigns the value 1 to x , then assigns the value 2 to y . While thread $t2$ prints the value of y , then the value of x . Both threads are allowed to finish before the program ends. The program contains two possible data races: thread $t1$ could be writing to x at the same time that thread $t2$ accesses x to print its value, and thread $t1$ could be writing to y at the same time that thread $t2$ accesses y to print its value. These two data races can be fixed by using mutexes to make the access of each variable mutually exclusive as shown below.

```
#include <iostream>
#include <thread>
#include <mutex>

int x = 0;
int y = 0;
std::mutex m_x;
std::mutex m_y;

int main()
{
    std::thread t1([ ] () {
        std::unique_lock lock_x(m_x, std::defer_lock);
        std::unique_lock lock_y(m_y, std::defer_lock);
        lock_x.lock();  x = 1;    lock_x.unlock();
        lock_y.lock();  y = 2;    lock_y.unlock();
    });
    std::thread t2([ ] () {
        std::unique_lock lock_x(m_x, std::defer_lock);
        std::unique_lock lock_y(m_y, std::defer_lock);
        lock_y.lock();  std::cout << y << " ";  lock_y.unlock();
        lock_x.lock();  std::cout << x << std::endl;  lock_x.unlock();
    });
    t1.join();
    t2.join();
}
```

Once these data races have been resolved, considering the SC-DRF memory model, there are six possible sequentially consistent orders that the program could follow, these are listed below along with the possible printed values of the variables. There is freedom for instructions A1 and B1 to occur concurrently. Also, instructions A2 and B2 can occur concurrently.

A1: $x = 1$, A2: $y = 2$, B1: $\text{std::cout} << y << \text{" "}$, B2: $\text{std::cout} << x << \text{std::endl}$

First instruction	Second instruction	Third instruction	Fourth instruction	Printed value of variable x at end of program	Printed value of variable y at end of program
A1	A2	B1	B2	1	2
A1	B1	A2	B2	1	0
A1	B1	B2	A2	1	0
B1	A1	A2	B2	1	0
B1	A1	B2	A2	1	0
B1	B2	A1	A2	0	0

7.1 b)

This program is similar to the previous, it creates two threads with shared variables x and y initially set to zero. Thread $t1$ assigns the value 1 to x , then assigns the value 2 to y . While thread $t2$ prints the value of y , then the value of x . Both threads are allowed to finish before the program ends. However, this program uses a single mutex to make the access to both variables mutually exclusive. This results in having no data races; however, this puts the entire code executed by each thread in a *happens-before* relationship. This causes only two possible sequentially consistent orders that the program could follow, as shown below. One case is that thread $t1$ locks the mutex first, executes its code, and unlocks the mutex; thread $t2$ can now lock the mutex, execute its code, and unlocks. The other case is vice-verse, thread $t2$ locks and executes first, then thread $t1$ follows. This forces the program to execute like it were a fully synchronous program.

A1: $x = 1$, A2: $y = 2$, B1: `std::cout << y << " "`, B2: `std::cout << x << std::endl`

First instruction	Second instruction	Third instruction	Fourth instruction	Printed value of variable x at end of program	Printed value of variable y at end of program
A1	A2	B1	B2	1	2
B1	B2	A1	A2	0	0

7.1 g)

This program creates two threads with shared variables x initially set to zero and *done* initially set to *false*. Thread $t1$ assigns 42 to x , then assigns *true* to *done*. Thread $t2$ actively reads the *done* and does not stop until *done* is *true*, then $t2$ uses an assertion to check if x is equals to 42. Both threads are allowed to finish before the program ends. With the way the program is written, there is only one possible data race that could occur: if $t1$ is writing to *done* while $t2$ is reading the state of *done*. This can be fixed by using a condition variable instead of actively checking *done* in $t2$. Due to the way the program is written, there are three possible sequentially consistent orders that the program could follow, as shown below. This is because B2 can not occur before A2, $t2$ will continue to execute B1 until A2 has been executed. Therefore, the assert in B2 will always be true.

A1: $x = 42$, A2: *done* = *true*, B1: `while (!done)`, B2: `assert(x == 42)`

First instruction	Second instruction	Third instruction	Fourth instruction
A1	A2	B1	B2
A1	B1	A2	B2
B1	A1	A2	B2

7.1 i)

This program creates two threads with shared variables x and y initially set to zero. Thread $t1$ checks if x is 1, and if it is, sets y to 1. Thread $t2$ checks if y is 1, and if it is, sets x to 1. Both threads are allowed to finish before the program ends. Due to the way the program is written, there is no data race. There are only two possible sequentially consistent orders that the program could follow:

1. Thread $t1$ checks the value of x in the if-statement which will result in *false*, thread $t2$ checks the value of y in the if-statement which will result in *false* as well.
2. The vice-versa occurs where $t2$ checks first, then $t1$.

7.1 l)

This program creates two threads with a shared variable *counter* initialized to zero. Each thread tries to increment *counter* 100,000 times. This will cause a data race when both threads try to write to *counter* at the same time. This can be fixed by making the access to *counter* mutually exclusive using a mutex. There is a wide number of ways in which the 100,000 incrementations of *counter* from each thread can be interleaved.

7.1 m)

This program creates two threads with a shared object w of class *Widget*. Objects of class *Widget* have two variables x and y , and two mutexes $xMutex$ and $yMutex$. Thread $t1$ locks $xMutex$ and sets x of object w to 1. Thread $t2$ locks $yMutex$ and sets y of object w to 1. There are no data races in this program. Since the threads are locking different mutexes, line 18 and line 24 can execute concurrently. In a sequentially consistent order, either line 18 could be executed first, then line 24, or vice-versa.

7.10)

A1: $x = 1;$

A2: $a = y;$

B1: $y = 1;$

B2: $b = x;$

All possible sequentially consistent executions of this program

First instruction	Second instruction	Third instruction	Fourth instruction	Value of variable a	Value of variable b
A1	A2	B1	B2	0	1
A1	B1	A2	B2	1	1
A1	B1	B2	A2	1	1
B1	A1	A2	B2	1	1
B1	A1	B2	A2	1	1
B1	B2	A1	A2	1	0

There is one combination of that cannot be obtained for a and b : where $a = 0$ and $b = 0$.

7.12 a)

In this scenario, the **assertion in instruction B2** will be **true sometimes**. The only two possible sequentially consistent executions that reaches the line of the assertion and makes it true are {A1,A2,B1,B2} and {A1,B1,A2,B2}, shown in the table below. Where A1 executing before B1 allows the program to reach the assertion line and A2 executing before B2 makes the assertion true. The execution order that reaches the assertion line and makes it false is {A1,B1,B2,A2}, where A1 executing before B1 allows the program to reach the assertion line and B2 executing before A2 makes the assertion false. Other sequentially consistent execution orders do not reach the assertion line because B1 is executed before A1.

First instruction	Second instruction	Third instruction	Fourth instruction	State of assertion in B2
A1	A2	B1	B2	True
A1	B1	A2	B2	True
A1	B1	B2	A2	False
B1	A1	A2	B2	Not reached
B1	A1	B2	A2	Not reached
B1	B2	A1	A2	Not reached

7.12 b)

In this scenario, if the program can reach the **assertion in instruction B2**, the assertion will **always be true**. This is because there is only one possible sequentially consistent execution order that can reach the assertion: {A1,A2,B1,B2}, this is because A2 is executed before B1. In this case, since A1 must be executed before A2 and B1 must be executed before B2, this results in A1 executing before B2, causing the assert to be true. Other sequentially consistent execution orders do not reach the assertion line because B1 is executed before A2. Also, it is not possible for B2 to execute before A2 because thread 2 actively executes B1 until thread 1 executes A2, therefore, B2 cannot be executed until thread 2 stops executing B1 which happens after thread 1 executes A2. This case is marked in **red**.

First instruction	Second instruction	Third instruction	Fourth instruction	State of assertion in B2
A1	A2	B1	B2	True
A1	B1	A2	B2	Not reached
A1	B1	B2	A2	Not reached
B1	A1	A2	B2	Not reached
B1	A1	B2	A2	Not reached
B1	B2	A1	A2	Not reached

7.12 c)

In this scenario, the program reaches both assertions only in four possible sequentially consistent execution orders as shown in below, where A2 is executed before B1. The **assertion in instruction B2 will always be true** if the program reaches that assertion. This is because in those four possibilities, A1 is executed before B2. The **assertion in instruction B3 will be true sometimes**. Only if A3 is executed before B3, then the assertion is true, this is the case in 3 out of the 4 possibilities. In the fourth, the assertion is false because B3 is executed before A3. Also, it is not possible for B2 or B3 to execute before A2 because thread 2 actively executes B1 until thread 1

executes A2, therefore, B2 or B3 cannot be executed until thread 2 stops executing B1 which happens after thread 1 executes A2. This case is marked in red.

First	Second	Third	Fourth	Fifth	Sixth	State of assertion in B2	State of assertion in B3
A1	A2	A3	B1	B2	B3	True	True
A1	A2	B1	A3	B2	B3	True	True
A1	A2	B1	B2	A3	B3	True	True
A1	A2	B1	B2	B3	A3	True	False
A1	B1	A2	A3	B2	B3	Not reached	Not reached
A1	B1	A2	B2	A3	B3		
A1	B1	A2	B2	B3	A3		
A1	B1	B2	A2	A3	B3		
A1	B1	B2	A2	B3	A3		
A1	B1	B2	B3	A2	A3		
B1	A1	A2	A3	B2	B3		
B1	A1	A2	B2	A3	B3		
B1	A1	A2	B2	B3	A3		
B1	A1	B2	A2	A3	B3		
B1	A1	B2	A2	B3	A3		
B1	A1	B2	B3	A2	A3		
B1	B2	A1	A2	A3	B3		
B1	B2	A1	A2	B3	A3		
B1	B2	A1	B3	A2	A3		
B1	B2	B3	A1	A2	A3		

Part C Report and Comment on the Julia set program

The average amount of time required for the execution of the *compute_julia_set* function is recorded below. The results obtained are to be expected because the system that the program ran on utilizes an intel core i7-8565u processor which has 4 cores. Therefore, it makes sense that the fastest computation will typically occur with 4 threads running where the system can delegate each core to run one thread if the system is able to. Also, as expected, running the program with multiple threads improves performance in terms of speed compared to running the program with a single thread.

Threads	Time (seconds)
1	0.973576
2	0.502833
4	0.305151
8	0.378237

```
1  # Specify Minimum Required Version
2  cmake_minimum_required(VERSION 3.1 FATAL_ERROR)
3
4  # Specify Project and Language
5  project(cpp_concurrency LANGUAGES CXX)
6
7  # Require compliance with C++17
8  set(CMAKE_CXX_STANDARD 17)
9  set(CMAKE_CXX_STANDARD_REQUIRED TRUE)
10
11 # Set Include Directory
12 include_directories(include)
13
14 # Find the threads library, indicating a preference for the pthread library
15 set(THREADS_PREFER_PTHREAD_FLAG true)
16 find_package(Threads REQUIRED)
17
18 # Add Executable Program
19 add_executable(test_queue app/test_queue.cpp)
20 add_executable(test_thread_pool app/test_thread_pool.cpp lib/thread_pool.cpp)
21 add_executable(test_julia_set app/test_julia_set.cpp lib/thread_pool.cpp)
22
23 # Set the libraries for the target
24 target_link_libraries(test_queue Threads::Threads)
25 target_link_libraries(test_thread_pool Threads::Threads)
26 target_link_libraries(test_julia_set Threads::Threads)
```

```
1  #ifndef QUEUEHPP
2  #define QUEUEHPP
3  #include <cstdint>
4  #include <queue>
5  #include <utility>
6  #include <mutex>
7  #include <thread>
8  #include <condition_variable>
9
10 namespace ra::concurrency {
11
12     // Concurrent bounded FIFO queue class.
13     template <class T>
14     class queue
15     {
16     public:
17
18         // The type of each of the elements stored in the queue.
19         using value_type = T;
20
21         // An unsigned integral type used to represent sizes.
22         using size_type = std::size_t;
23
24         // A type for the status of a queue operation.
25         enum class status {
26             success = 0, // operation successful
27             empty, // queue is empty (not currently used)
28             full, // queue is full (not currently used)
29             closed, // queue is closed
30         };
31
32         // A queue is not default constructible.
33         queue() = delete;
34
35         // Constructs a queue with a maximum size of max_size.
36         // The queue is marked as open (i.e., not closed).
37         // Precondition: The quantity max_size must be greater than
38         // zero.
39         queue(size_type max_size) : capacity_(max_size), stat_(status::empty)
40     ) {}
41
42     // A queue is not movable or copyable.
43     queue(const queue&) = delete;
44     queue& operator=(const queue&) = delete;
45     queue(queue&&) = delete;
46     queue& operator=(queue&&) = delete;
47
48     // Destroys the queue after closing the queue (if not already
49     // closed) and clearing the queue (if not already empty).
50     ~queue() {
51         close();
52         clear();
53     }
54
55     // Inserts the value x at the end of the queue, blocking if
56     // necessary.
57     // If the queue is full, the thread will be blocked until the
58     // queue insertion can be completed or the queue is closed.
59     // If the value x is successfully inserted on the queue, the
60     // function returns status::success.
61     // If the value x cannot be inserted on the queue (due to the
62     // queue being closed), the function returns with a return
```

```
62      // value of status::closed.
63      // This function is thread safe.
64      // Note: The rvalue reference parameter is intentional and
65      // implies that the push function is permitted to change
66      // the value of x (e.g., by moving from x).
67      status push(value_type&& x){
68          std::unique_lock<std::mutex> lock(m_);
69          c_full_.wait(lock, [this]() { return (stat_!= status::full); });
70          if(stat_==status::closed){
71              return stat_;
72          }
73          else {
74              queue_.push(std::move(x));
75              if(queue_.size()==capacity_){
76                  stat_ = status::full;
77              }
78              if(stat_==status::empty){
79                  stat_ = status::success;
80              }
81              c_empty_.notify_one();
82              return (status::success);
83          }
84          //lock.unlock();
85          //c_empty_.notify_one();
86          //return (status::success);
87      }
88
89      // Removes the value from the front of the queue and places it
90      // in x, blocking if necessary.
91      // If the queue is empty and not closed, the thread is blocked
92      // until: 1) a value can be removed from the queue; or 2) the
93      // queue is closed.
94      // If the queue is closed, the function does not block and either
95      // returns status::closed or status::success, depending on whether
96      // a value can be successfully removed from the queue.
97      // If a value is successfully removed from the queue, the value
98      // is placed in x and the function returns status::success.
99      // If a value cannot be successfully removed from the queue (due to
100      // the queue being both empty and closed), the function returns
101      // status::closed.
102      // This function is thread safe.
103      status pop(value_type& x){
104          std::unique_lock<std::mutex> lock(m_);
105          c_empty_.wait(lock, [this]() { return (stat_!= status::empty); });
106          if((stat_==status::closed) && (queue_.empty())){
107              return stat_;
108          }
109          else{
110              x = queue_.front();
111              queue_.pop();
112              if(stat_==status::full){
113                  stat_ = status::success;
114              }
115              if((stat_!=status::closed) && (queue_.empty())){
116                  stat_ = status::empty;
117              }
118              c_full_.notify_one();
119              return (status::success);
120          }
121          //lock.unlock();
122          //c_full_.notify_one();
123          //return (status::success);
```

```
124     }
125
126     // Closes the queue.
127     // The queue is placed in the closed state.
128     // The closed state prevents more items from being inserted
129     // on the queue, but it does not clear the items that are
130     // already on the queue.
131     // Invoking this function on a closed queue has no effect.
132     // This function is thread safe.
133     void close() {
134         std::scoped_lock<std::mutex> lock(m_);
135         if(stat_ != status::closed) {
136             stat_ = status::closed;
137         }
138         c_full_.notify_all();
139         c_empty_.notify_all();
140     }
141
142     // Clears the queue.
143     // All of the elements on the queue are discarded.
144     // This function is thread safe.
145     void clear() {
146         std::scoped_lock<std::mutex> lock(m_);
147         if(stat_ != status::empty || (!queue_.empty())) {
148             while(! (queue_.empty())) {
149                 queue_.pop();
150             }
151             if(stat_ != status::closed) {
152                 stat_ = status::empty;
153                 c_full_.notify_all();
154             }
155         }
156     }
157
158     // Returns if the queue is currently full (i.e., the number of
159     elements in the queue equals the maximum queue size).
160     // This function is not thread safe.
161     bool is_full() const {
162         //std::scoped_lock<std::mutex> lock(m_);
163         return (queue_.size()==capacity_);
164     }
165
166     // Returns if the queue is currently empty.
167     // This function is not thread safe.
168     bool is_empty() const {
169         //std::scoped_lock<std::mutex> lock(m_);
170         return (queue_.empty());
171     }
172
173     // Returns if the queue is closed (i.e., in the closed state).
174     // This function is not thread safe.
175     bool is_closed() const {
176         //std::scoped_lock<std::mutex> lock(m_);
177         return (stat_==status::closed);
178     }
179
180     // Returns the maximum number of elements that can be held in
181     the queue.
182     // This function is not thread safe.
183     size_type max_size() const {
184         //std::scoped_lock<std::mutex> lock(m_);
185         return capacity_;
```

```
186         }
187
188     private:
189         std::queue<value_type> queue_;
190         mutable std::mutex m_;
191         mutable std::condition_variable c_empty_;
192         mutable std::condition_variable c_full_;
193
194         size_type capacity_;
195         status stat_;
196
197     };
198
199
200
201
202 }
203 #endif
```

```
1  #include "ra/queue.hpp"
2  #include <cstdint>
3  #include <thread>
4  #include <mutex>
5  #include <iostream>
6  #include <utility>
7
8  using value_type = double;
9  using queue = ra::concurrency::queue<value_type>;
10 using status = ra::concurrency::queue<value_type>::status;
11 using queueForStat = ra::concurrency::queue<status>;
12 using std::cout;
13 using std::endl;
14
15 constexpr std::size_t num_elements(10);
16 queue q(2); //std::mutex m_q;
17 queue result(10);
18 queueForStat stat_store(30);
19
20 int main(){
21     std::thread t1([](){
22         //std::unique_lock<std::mutex> lock(m_q, std::defer_lock);
23         for(std::size_t i=0; i<num_elements; ++i){
24             //std::unique_lock<std::mutex> lock(m_q);
25             status stat = q.push(value_type(i)); // lock.unlock();
26             stat_store.push(std::move(stat));
27         }
28         q.close();
29     });
30     std::thread t2([](){
31         bool keep_running(true);
32         //lstd::unique_lock<std::mutex> lock(m_q, std::defer_lock);
33         //while(keep_running){
34         for(std::size_t i=0; i<num_elements; ++i){
35             //lstd::unique_lock<std::mutex> lock(m_q);
36             //if((q.is_closed()) && (q.is_empty())){ keep_running = false; } // 1
37             ock.unlock();
38             value_type x;
39             status stat = q.pop(x);
40             stat_store.push(std::move(stat));
41             x *= value_type(2);
42             stat = result.push(std::move(x));
43             stat_store.push(std::move(stat));
44             //if((q.is_closed()) && (q.is_empty())){ keep_running = false; } // 1
45             ock.unlock();
46         }
47     });
48     t1.join(); t2.join();
49     cout<<"Is q closed (should be true): "<<(q.is_closed())<<endl;
50     value_type res(0);
51     // Print out the results
52     while(!result.is_empty()){
53         result.pop(res);
54         cout<<res<<" ";
55     }
56     cout<<endl;
57     cout<<"Is stat_store full (should be true): "<<(stat_store.is_full())<<endl;
58     cout<<"Max size of stat_store: "<<(stat_store.max_size())<<endl;
59     //stat_store.clear();
60     // Print out the status
61     status stat(status::success);
```

```
61     while(!stat_store.is_empty()){
62         stat_store.pop(stat);
63         if(stat==status::success){ cout<<"success"<<endl; }
64         else if(stat==status::closed){ cout<<"closed"<<endl; }
65         else { cout<<"Problem here"<<endl; }
66     }
67
68 }
```



```

1  #ifndef THREADPOOLHPP
2  #define THREADPOOLHPP
3  #include "ra/queue.hpp"
4  #include <cstdint>
5  #include <vector>
6  #include <mutex>
7  #include <thread>
8  #include <condition_variable>
9  #include <functional>
10 #include <utility>
11
12
13 namespace ra::concurrency {
14
15     // Thread pool class.
16     class thread_pool
17     {
18     public:
19
20         // An unsigned integral type used to represent sizes.
21         using size_type = std::size_t;
22
23         using func = std::function<void()>;
24         using my_queue = typename ra::concurrency::queue<func>;
25
26         // Function performed by each thread
27         friend void worker(thread_pool*);
28         /*void worker(){
29             std::unique_lock<std::mutex> lock(m_);
30             while(!(shutDown_ && queue_.is_empty())){
31                 c_task_.wait(lock, [this]() { return ((allThreadsMade_ && (!queue_
32 _.is_empty())) || (shutDown_));});
33                 if(!queue_.is_empty()){
34                     func task_;
35                     queue_.pop(task_);
36                     lock.unlock();
37                     c_add_.notify_one();
38                     task_();
39                     lock.lock();
40                 }
41                 c_done_.notify_all();
42                 shutDownFinished_ = true;
43             */
44
45         // Creates a thread pool with the number of threads equal to the
46         // hardware concurrency level (if known); otherwise the number of
47         // threads is set to 2.
48         thread_pool() : allThreadsMade_(false), shutDown_(false), shutDownFinish
49 ed_(false), queue_(64) {
50             size_type num_threads = std::thread::hardware_concurrency();
51             if(num_threads == size_type(0)) { num_threads = size_type(2); }
52             num_threads_ = num_threads;
53             void worker(thread_pool*);
54             for(size_type i=0; i<num_threads; ++i){
55                 workers.emplace_back(worker, this);
56             }
57             std::scoped_lock<std::mutex> lock(m_);
58             allThreadsMade_ = true;
59             c_task_.notify_one();
60         }

```

```

61         // Creates a thread pool with num_threads threads.
62         // Precondition: num_threads > 0
63         thread_pool(std::size_t num_threads) : allThreadsMade_(false), shutDown_
        (false), shutDownFinished_(false), queue_(64), num_threads_(num_threads) {
64             void worker(thread_pool*);
65             for(size_type i=0; i<num_threads; ++i){
66                 workers.emplace_back(worker, this);
67             }
68             std::scoped_lock<std::mutex> lock(m_);
69             allThreadsMade_ = true;
70             c_task_.notify_one();
71         }
72
73         // A thread pool is not copyable or movable.
74         thread_pool(const thread_pool&) = delete;
75         thread_pool& operator=(const thread_pool&) = delete;
76         thread_pool(thread_pool&&) = delete;
77         thread_pool& operator=(thread_pool&&) = delete;
78
79         // Destroys a thread pool, shutting down the thread pool first
80         // (if not already shutdown).
81         ~thread_pool(){
82             shutdown();
83             for(auto& i : workers){
84                 if(i.joinable()){
85                     i.join();
86                 }
87             }
88         }
89
90         // Gets the number of threads in the thread pool.
91         // This function is not thread safe.
92         size_type size() const{
93             return num_threads_;
94         }
95
96         // Enqueues a task for execution by the thread pool.
97         // This function inserts the task specified by the callable
98         // entity func into the queue of tasks associated with the
99         // thread pool.
100        // This function may block if the number of currently
101        // queued tasks is sufficiently large.
102        // Note: The rvalue reference parameter is intentional and
103        // implies that the schedule function is permitted to change
104        // the value of func (e.g., by moving from func).
105        // Precondition: The thread pool is not in the shutdown state
106        // and is not currently in the process of being shutdown via
107        // the shutdown member function.
108        // This function is thread safe.
109        void schedule(std::function<void()>&& func){
110            std::unique_lock<std::mutex> lock(m_);
111            /*if(shutDown_){
112                c_done_.wait(lock, [this]() { return (queue_.is_empty()); });
113            }
114            else{*/
115            c_add_.wait(lock, [this]() { return ((!queue_.is_full()) || (shut
Down_)); });
116            if(!shutDown_){
117                queue_.push(std::move(func));
118            }
119            //}
120        }

```

```
121
122     // Shuts down the thread pool.
123     // This function places the thread pool into a state where
124     // new tasks will no longer be accepted via the schedule
125     // member function.
126     // Then, the function blocks until all queued tasks
127     // have been executed and all threads in the thread pool
128     // are idle (i.e., not currently executing a task).
129     // Finally, the thread pool is placed in the shutdown state.
130     // If the thread pool is already shutdown at the time that this
131     // function is called, this function has no effect.
132     // After the thread pool is shutdown, it can only be destroyed.
133     // This function is thread safe.
134     void shutdown() {
135         std::unique_lock<std::mutex> lock(m_);
136         if(!shutDown_){
137             shutDown_ = true;
138             queue_.close();
139             c_task_.notify_all();
140             c_done_.wait(lock, [this]() { return (queue_.is_empty()); });
141         }
142     }
143
144     // Tests if the thread pool has been shutdown.
145     // This function is not thread safe.
146     bool is_shutdown() const{
147         return shutDownFinished_;
148     }
149
150     private:
151     std::vector<std::thread> workers;
152     mutable std::mutex m_;
153     mutable std::condition_variable c_task_;
154     mutable std::condition_variable c_done_;
155     mutable std::condition_variable c_add_;
156
157     my_queue queue_;
158     bool allThreadsMade_;
159     bool shutDown_;
160     bool shutDownFinished_;
161     size_type num_threads_;
162
163 };
164
165
166
167
168 }
169 #endif
```

```
1  #include "ra/thread_pool.hpp"
2
3  namespace ra::concurrency {
4      void worker(thread_pool* obj){
5          std::unique_lock<std::mutex> lock(obj->m_);
6          while(!(obj->shutDown_ && obj->queue_.is_empty())){
7              obj->c_task_.wait(lock, [obj](){ return ((obj->allThreadsMade_ && (!
obj->queue_.is_empty())) || (obj->shutDown_));});
8              if(!obj->queue_.is_empty()){
9                  std::function<void()> task_;
10                 obj->queue_.pop(task_);
11                 lock.unlock();
12                 obj->c_add_.notify_one();
13                 task_();
14                 lock.lock();
15             }
16         }
17         obj->shutDownFinished_ = true;
18         obj->c_done_.notify_all();
19     }
20
21 }
```

```
1  #include "ra/thread_pool.hpp"
2  #include <iostream>
3  #include <cstdint>
4
5  using thr_pool = ra::concurrency::thread_pool;
6  using std::cout;
7  using std::endl;
8
9  int main(){
10     thr_pool tp(8);
11     for(std::size_t i=0; i<30; ++i){
12         tp.schedule([]() {cout<<"hello world"<<endl;});
13     }
14     tp.shutdown();
15     //cout<<"Is shutdown (must be false): "<<(tp.is_shutdown())<<endl;
16     //cout<<"Num of threads "<<(tp.size())<<endl;
17     //while(!tp.is_shutdown()){
18     //cout<<"Is shutdown (must be true): "<<(tp.is_shutdown())<<endl;
19
20 }
```

```

1  #ifndef JULIASETHPP
2  #define JULIASETHPP
3  #include "ra/thread_pool.hpp"
4  #include <complex>
5  #include <boost/multi_array.hpp>
6  #include <cstdint>
7  #include <thread>
8  #include <mutex>
9  #include <functional>
10
11
12 namespace ra::fractal {
13     using size_type = int;
14
15     boost::multi_array<int, 2> a_; // shared pointer to matrix a
16     std::mutex m_a; // Mutex used to make access to a_ptr mutually exclusive
17
18     template <class Real>
19     std::complex<Real> getZ(size_type l, size_type k, size_type W, size_type H,
20 std::complex<Real> U, std::complex<Real> V) {
21         Real U0 = U.real();
22         Real U1 = U.imag();
23         Real V0 = V.real();
24         Real V1 = V.imag();
25         Real Z0 = U0 + ( (Real(k)/Real(W-1)) * (V0-U0) );
26         Real Z1 = U1 + ( (Real(l)/Real(H-1)) * (V1-U1) );
27         std::complex<Real> Z(Z0,Z1);
28         return Z;
29     }
30
31     template <class Real>
32     int julia_ym(int m, std::complex<Real> z, std::complex<Real> c){
33         int i(0);
34         while(!( (i>=m) || (std::abs(z)>Real(2)) )){
35             z = (z * z) + c;
36             ++i;
37         }
38         if(i>=m) { return m; }
39         else{ return i; }
40     }
41
42     template <class Real>
43     void compute_julia_set(const std::complex<Real>& bottom_left, const std::com
44 plex<Real>& top_right, const std::complex<Real>& c, int max_iters, boost::multi_
45 array<int, 2>& a, int num_threads){
46         size_type W = a.shape()[1]; // num of cols
47         size_type H = a.shape()[0]; // num of rows
48         a_.resize(boost::extents[a.shape()[0]][a.shape()[1]]);
49         a_ = a;
50         //a_ = a;
51         ra::concurrency::thread_pool tp(num_threads); // Create Thread pool
52         for(size_type l=0; l<H; ++l){
53             //judeboost::multi_array<int, 2>* a_row_ptr = &(a[l][0]);
54             tp.schedule([l, W, H, bottom_left, top_right, c, max_iters](){
55                 std::unique_lock lock(m_a, std::defer_lock);
56                 for(size_type k=0; k<W; ++k){
57                     std::complex<Real> z = getZ(l,k,W,H,bottom_left, top_right);
58                     int a_ym = julia_ym(max_iters,z,c);
59                     lock.lock();
60                     a_[l][k] = a_ym;
61                     lock.unlock();

```

```
60         }
61         });
62     }
63     tp.shutdown();
64     std::unique_lock lock(m_a);
65     a.resize(boost::extents[a_.shape()[0]][a_.shape()[1]]);
66     a = a_;
67 }
68
69
70
71
72
73
74 }
75 #endif
```

```
1  #include "ra/julia_set.hpp"
2  #include <complex>
3  #include <boost/multi_array.hpp>
4  #include <cstdint>
5  #include <iostream>
6  #include <fstream>
7  #include <chrono>
8
9  int main(){
10     using Real = double;
11     using std::cout;
12     using std::endl;
13     std::complex<Real> bottom_left(-1.25,-1.25);
14     std::complex<Real> top_right(1.25,1.25);
15     std::complex<Real> c(0.37,-0.16);
16     int max_iters(255);
17     int num_threads(4);
18     int num_rows(512);
19     int num_cols(512);
20     boost::multi_array<int, 2> a(boost::extents[num_rows][num_cols]);
21     auto start = std::chrono::steady_clock::now();
22     ra::fractal::compute_julia_set(bottom_left,top_right,c,max_iters,a,num_threa
ds);
23     auto end = std::chrono::steady_clock::now();
24     std::chrono::duration<Real> elapsed_time = end - start;
25     cout<<"For "<<num_threads<<" threads, elapsed time: "<<(elapsed_time.count())<<" seconds
"<<endl;
26
27     std::ofstream outFile("image.pnm");
28     outFile<<"P2 "<<num_cols<<" "<<num_rows<<" "<<max_iters<<"\n";
29     for(int row(num_rows-1); row>=0; --row){
30         for(int col(0); col<num_cols; ++col){
31             if(col){ outFile<<" "; }
32             outFile<<(a[row][col]);
33         }
34         outFile<<"\n";
35     }
36     outFile.close();
37
38
39 }
```